Human-in-the-Loop Design Checklist

Right-sized human oversight for AI agents — designed one action at a time.

How to use it. Skip the unhelpful yes/no question “should there be a human in the loop?” Decide it **per action** with four moves: **Grade, Guard, Show, Prove.** Work each action through all four.

GRADE every action (G0–G3)

Score each action on its own — blast radius, reversibility, and severity together set its grade.

- [ ] **Blast radius.** Have you scored how far harm could spread if this action goes wrong?
- [ ] **Reversibility.** Can the action be undone, and how fast?
- [ ] **Severity.** How bad is the worst outcome if it is wrong?
- [ ] **Grade assigned.** Does each action carry a grade from G0 (trivial) to G3 (critical), set from blast radius, reversibility, and severity together?

GUARD by grade

Match the oversight to the grade — and prevent by design where review cannot keep up.

- [ ] **Oversight matched to grade.** Do low grades log-and-go while high grades get real review?
- [ ] **Prevent where review cannot keep up.** Where a human cannot reliably catch the mistake in time, do you prevent the bad outcome by design (sandbox it, make it reversible, or block it) instead of asking for approval?
- [ ] **No approval theater.** Can every approval step actually stop the action before it takes effect, not just record a name after?
- [ ] **Irreversible, high-severity actions are gated or prevented,** never left to a rubber stamp.

SHOW the human a good oversight moment

When you do ask a person, make the moment one they can actually judge.

- [ ] **Clear action and consequence.** Does the prompt make what is about to happen, and its impact, obvious?
- [ ] **Provenance.** Can the reviewer see how the request got to them — the chain that produced it?
- [ ] **Decision support.** Are the likely problems surfaced rather than buried?
- [ ] **Right-sized attention.** Does the oversight cost no more of the person’s attention than the action is worth?

PROVE it works

An oversight loop you have not tested is a loop you cannot trust.

- [ ] **Test the oversight.** Have you run real mistakes through the loop to confirm a human actually catches them?
- [ ] **Measure the catch rate.** Do you track how often oversight catches a bad action versus waves it through?
- [ ] **Re-test on change.** When the agent, model, or task changes, do you re-prove the loop still catches mistakes?

Put it on the rails. Grade every action, guard it to its grade, show the human a moment they can actually judge, and prove the loop catches mistakes before you trust it. **Re-grade when the agent or its task changes.**