

YOUR AI DEMO WORKS.

PRODUCTION HAS FOLLOW-UP QUESTIONS.

The agent gives a convincing answer. But did the database change? Did the citation support the claim? Did the patch solve the problem, or merely charm the test?

Measure Twice, Prompt Once shows product and engineering teams how to turn those awkward questions into a working evaluation loop.

INSIDE

- ◆ Read real failures without drowning in traces
- ◆ Build tasks and graders around evidence that matters
- ◆ Compare changes without statistical theater
- ◆ Connect CI, release gates, and production feedback
- ◆ Give the loop an owner, a cadence, and a budget


Four published programs. One complete fictional loop. Plus statistical recipes, thirteen exercises, sixteen templates, and a 30-day minimum viable eval-loop plan.

ABOUT THE AUTHOR

Brenn Hill is a software engineer, engineering leader, and published author. He teaches AI-augmented development practices and builds tools for verification, agent loops, and human oversight.

MEASURE TWICE, PROMPT ONCE

Building Evals for AI Agents That Work in Production



BRENN HILL

MEASURE TWICE, PROMPT ONCE

◆

BRENN HILL